


Lecture 25: Deep neural networks continued

Professor Ilias Bilonis


Image classification using convolutional neural networks

Animal visual cortex

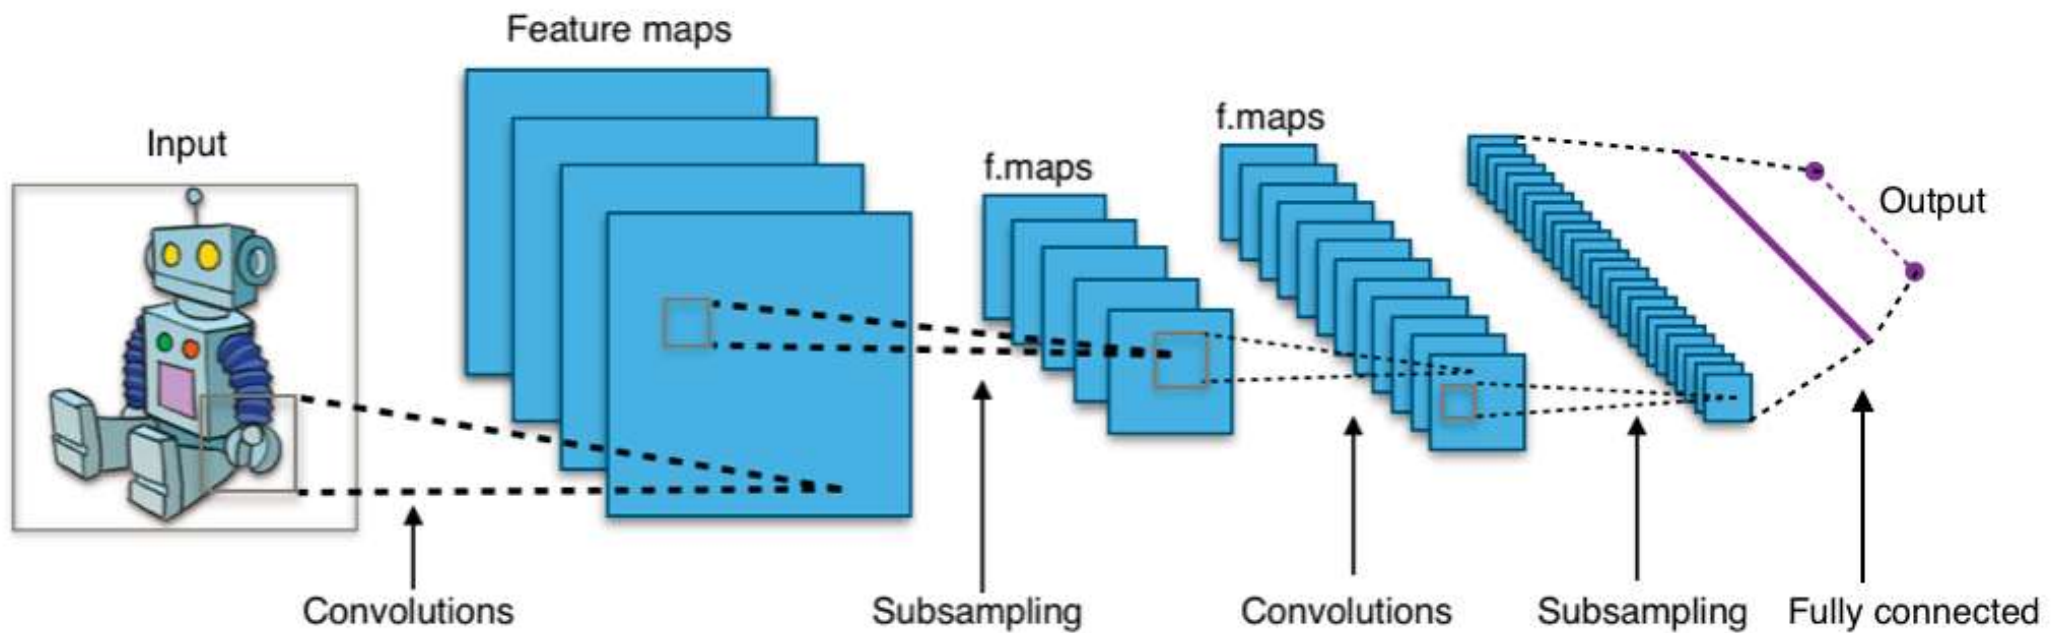


https://en.wikipedia.org/wiki/Visual_cortex#/media/File:Neural_pathway_diagram.svg

https://en.wikipedia.org/wiki/Convolutional_neural_network#/media/File:Conv_layer.png



A typical CNN architecture




https://en.wikipedia.org/wiki/Convolutional_neural_network#/media/File:Typical_cnn.png


Why do they work well for image recognition?

- Translation invariant.
- Shared parameters -> Less prone to overfitting.
- Hierarchical pattern recognition.

The convolutional layer




The pooling layer (max)



Example of pooling

Max Pooling
 $D=4$ $P=2$ $\Rightarrow d = \frac{D}{P} = 2$



https://en.wikipedia.org/wiki/Convolutional_neural_network#Receptive_fields_in_the_visual_cortex